

SAMSUNG

Voice recognition

Before and after Deep Learning

Mikhail Kudinov

Samsung R&D Center Russia

March 29, 2019

Table of Contents

Introduction

Intro subsection

Section 1

Section 2

Conclusion

Acknowledgements

We are here!

Introduction

Intro subsection

Section 1

Section 2

Conclusion

Acknowledgements

This is the first slide


Ideas

See how this looks


- ▶ Item 1
- ▶ Item 2

Universal Background Model

- ▶ Criterion:
 $\Lambda(X) = \log p(X|\theta_s) - \log p(X|\theta_{bck}) \geq C$
- ▶ Train separate model θ_{bck} called Universal Background Model
- ▶ All frames are assumed to be independent
- ▶ $p(X_t|\theta)$ is a GMM with diagonal covariance matrices Σ_i
- ▶ θ_{bck} is trained with EM algorithm on the whole collection of speech data
- ▶ Speaker model θ_s is an MAP update of θ_{bck}



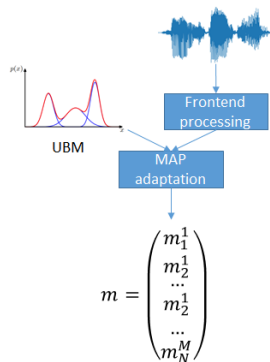
Likelihood ratio based system layout



MAP update of UBM

SVM on GMM supervectors

- ▶ Supervector m : do MAP-updates of means of the UBM and concatenate them
- ▶ Use lower-bound on KL-divergence as a distance measure:
$$d(m^a, m^b) = \frac{1}{2} \sum_i \pi_i (m_i^a - m_i^b) \Sigma_i^{-1} (m_i^a - m_i^b)$$
- ▶ Use SVM on supervectors with the kernel defined above
- ▶ Nuisance attribute projection



Supervector computation

Extensions of supervector model

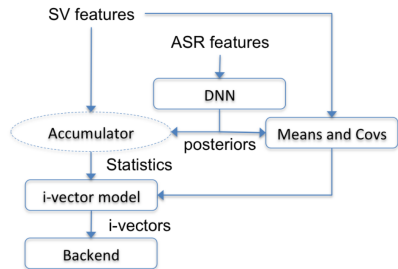
- ▶ SVM with simple cosine kernel on preprocessed supervectors
- ▶ Within-class Covariance Normalization $k(m_1, m_2) = m_1 W^{-1} m_2$,
where W is a covariance matrix: $W = \frac{1}{S} \sum_{s=1}^S \frac{1}{n_s} \sum_{i=1}^{n_s} (m_i^s - \overline{m_s})(m_i^s - \overline{m_s})^t$
- ▶ SVM on LDA-projected supervectors
- ▶ Nuisance Attribute Projection $\arg \min_P \sum_{i,j} W_{i,j} \|Pm_1 - Pm_2\|_2^2$, where $W_{i,j} = 0$ iff w_1 and w_2 are from e.g. different microphones (nuisance channel direction)
- ▶ Joint Factor Analysis: $m = \mu + Vy + Ux + Dz$, where y is a speaker-dependent and x is a channel-dependent vector
- ▶ Probabilistic LDA (PLDA)

i-vectors

- ▶ Front-end Factor Analysis for a supervector M : $m = \mu + Tw$, where μ is the speaker- and channel-independent vector (UBM-vector)
- ▶ T is a rectangular matrix of low rank
- ▶ $w \sim \mathcal{N}(0, \mathbb{I})$ is called *i-vector*
- ▶ i-vector estimation for known matrix T for utterance u :
 - ▶ Use UBM Ω to get Viterbi or Baum-Welch estimations for GMM component c :
 $N_c = \sum_t P(c|y_t, \Omega)$; $F_c = \sum_t P(c|y_t, \Omega)(y_t - \mu_c)$
 - ▶ $w = (I + T^T \Sigma^{-1} N T)^{-1} T^T \Sigma^{-1} F$, where N and F are concatenations of N_c and F_c for all GMM components $\{c_i\}$
- ▶ T is estimated using ML algorithm

DNNs for better i-vector extraction

- ▶ Use DNN to obtain class-posterior for statistics collection N and F for i-vector calculation
- ▶ Classes are context-aware *senones*
- ▶ Requires ASR system trained separately
- ▶ Similarity score was computed on PLDA-projections



DNN/i-vector framework

```
OLTP test statistics:
queries performed:
  read:                716114
  write:               255715
  other:               102290
  total:               1074119
transactions:         51139 (852.18 per sec.)
deadlocks:            12 (0.20 per sec.)
read/write requests:  971829 (16194.50 per sec.)
other operations:     102290 (1704.55 per sec.)
```

```
General statistics:
total time:           60.0098s
total number of events: 51139
total time taken by event execution: 479.6358
response time:
  min:                2.03ms
  avg:                 9.38ms
  max:                77.61ms
  approx. 95 percentile: 14.18ms
```

```
Threads fairness:
events (avg/stddev):   6392.3750/15.62
execution time (avg/stddev): 59.9545/0.00
```

```
OLTP test statistics:
queries performed:
  read:                747292
  write:               266853
  other:               106745
  total:               1120890
transactions:         53367 (889.03 per sec.)
deadlocks:            11 (0.18 per sec.)
read/write requests:  1014145 (16894.47 per sec.)
other operations:     106745 (1778.25 per sec.)
```

```
General statistics:
total time:           60.0282s
total number of events: 53367
total time taken by event execution: 479.6704
response time:
  min:                1.93ms
  avg:                 8.99ms
  max:                101.82ms
  approx. 95 percentile: 13.52ms
```

```
Threads fairness:
events (avg/stddev):   6670.8750/11.55
execution time (avg/stddev): 59.9588/0.01
```

This slide has code blocks

```
1      import numpy as np
2
3      def incmatrix(genl1,genl2):
4          m = len(genl1)
5          n = len(genl2)
6          M = None #to become the incidence matrix
7          VT = np.zeros((n*m,1), int) #dummy variable
```

We are here!

Introduction

Intro subsection

Section 1

Section 2

Conclusion

Acknowledgements

This is the second slide

A bit more information about this

Alert block

Alert text

We are here!

Introduction

Intro subsection

Section 1

Section 2

Conclusion

Acknowledgements

Generic slide

A bit more information about this

We are here!

Introduction

Intro subsection

Section 1

Section 2

Conclusion

Acknowledgements

Generic slide

A bit more information about this

We are here!

Introduction

Intro subsection

Section 1

Section 2

Conclusion

Acknowledgements

Generic slide

A bit more information about this

We are here!

Introduction

Intro subsection

Section 1

Section 2

Conclusion

Acknowledgements

Generic slide

A bit more information about this